

50.020 Network Security

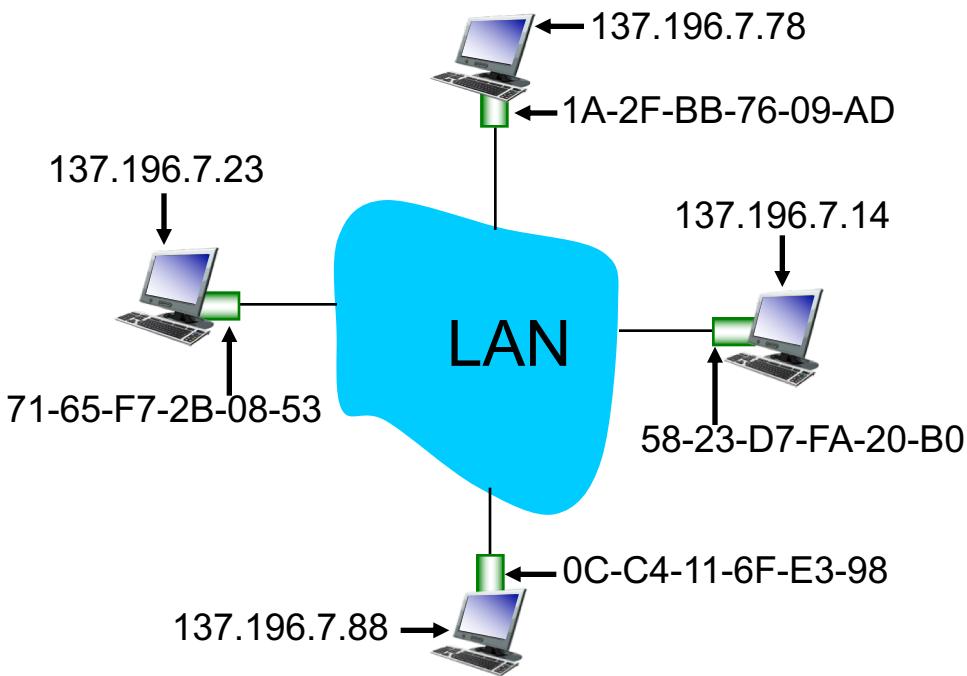
Lab 0: ARP sniffing and spoofing

Summary of last lecture

- Definition of Internet
- Circuit switching vs. packet switching
- Encapsulation and layered network model
- Security issues in computer networking
- Review ARP protocol

ARP: address resolution protocol

Question: how to determine interface's MAC address, knowing its IP address?



ARP table: each IP node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:
< IP address; MAC address; TTL >
- TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)
- Try : **arp -a**

ARP protocol: same LAN

- A wants to send datagram to B
 - B's MAC address not in A's ARP table.
- A **broadcasts** ARP query packet, containing B's IP address
 - destination MAC address = FF-FF-FF-FF-FF-FF
 - all nodes on LAN receive ARP query
- B receives ARP packet, replies to A with its (B's) MAC address
 - frame sent to A's MAC address (unicast)
- A caches (saves) IP-to-MAC address pair in its ARP table until information becomes old (times out)
 - soft state: information that times out (goes away) unless refreshed
- ARP is “plug-and-play”:
 - nodes create their ARP tables *without intervention from net administrator*

Example of an ARP request

The screenshot shows a Wireshark capture of an ARP request. The packet list view shows four ARP frames. The first three are from a host named 'HewlettP_bf:91:ee' with source MAC address 00:25:b3:bf:91:ee and destination MAC address Dell_c0:56:f0. These frames have lengths of 60 bytes and are ARP requests for the IP address 172.16.0.107. The fourth frame is a broadcast ARP request from the same host to all hosts (Broadcast) with length 60 bytes. The details view for the first frame shows the Ethernet II header, the ARP header (Type: ARP, Opcode: request), and the ARP payload. The payload shows the sender MAC as HewlettP_bf:91:ee, sender IP as 172.16.0.1, target MAC as Dell_c0:56:f0, and target IP as 172.16.0.107.

No.	Time	Source	Destination	Protocol	Length	Info
54	4.646389	HewlettP_bf:91:ee	Dell_c0:56:f0	ARP	60	Who has 172.16.0.107? Tell 172.16.0.1
55	4.646442	Dell_c0:56:f0	HewlettP_bf:91:ee	ARP	42	172.16.0.107 is at 00:21:70:c0:56:f0
56	4.646455	HewlettP_bf:91:ee	Dell_c0:56:f0	ARP	60	172.16.0.1 is at 00:25:b3:bf:91:ee
165	14.392559	HewlettP_bf:91:ee	Broadcast	ARP	60	Who has 172.16.0.1? Tell 172.16.0.105

Frame 54: 60 bytes on wire (480 bits), 60 bytes captured (480 bits)
Ethernet II, Src: HewlettP_bf:91:ee (00:25:b3:bf:91:ee), Dst: Dell_c0:56:f0 (00:21:70:c0:56:f0)
 ► Destination: Dell_c0:56:f0 (00:21:70:c0:56:f0)
 ► Source: HewlettP_bf:91:ee (00:25:b3:bf:91:ee)
 Type: ARP (0x0806)
 Padding: 00
Address Resolution Protocol (request)
 Hardware type: Ethernet (1)
 Protocol type: IPv4 (0x0800)
 Hardware size: 6
 Protocol size: 4
 Opcode: request (1)
 Sender MAC address: HewlettP_bf:91:ee (00:25:b3:bf:91:ee)
 Sender IP address: 172.16.0.1
 Target MAC address: 00:00:00_00:00:00 (00:00:00:00:00:00)
 Target IP address: 172.16.0.107

Example of an ARP reply

GENERAL_arppoison_chrianders.org_.pcap

arp

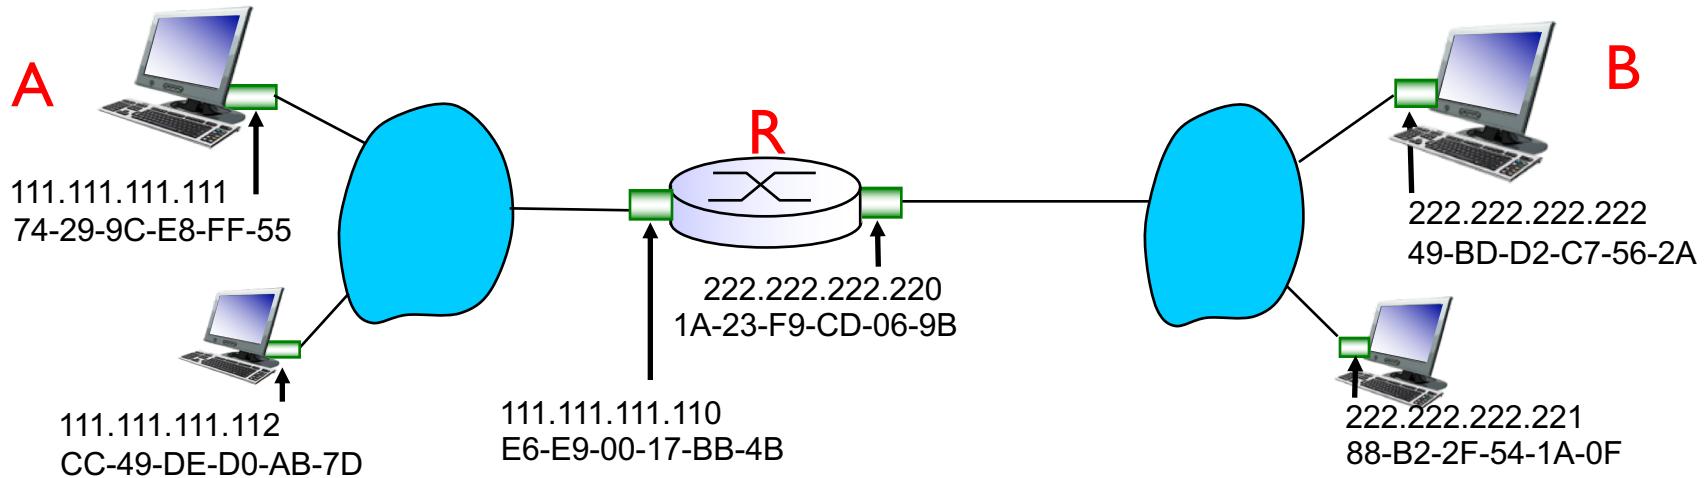
No.	Time	Source	Destination	Protocol	Length	Info
54	4.646389	HewlettP_bf:91:ee	Dell_c0:56:f0	ARP	60	Who has 172.16.0.107? Tell 172.16.0.1
55	4.646442	Dell_c0:56:f0	HewlettP_bf:91:ee	ARP	42	172.16.0.107 is at 00:21:70:c0:56:f0
56	4.646455	HewlettP_bf:91:ee	Dell_c0:56:f0	ARP	60	172.16.0.1 is at 00:25:b3:bf:91:ee
165	14.392559	HewlettP_bf:91:ee	Broadcast	ARP	60	Who has 172.16.0.1? Tell 172.16.0.105

▶ Frame 55: 42 bytes on wire (336 bits), 42 bytes captured (336 bits)
▼ Ethernet II, Src: Dell_c0:56:f0 (00:21:70:c0:56:f0), Dst: HewlettP_bf:91:ee (00:25:b3:bf:91:ee)
 ▶ Destination: HewlettP_bf:91:ee (00:25:b3:bf:91:ee)
 ▶ Source: Dell_c0:56:f0 (00:21:70:c0:56:f0)
 Type: ARP (0x0806)
▼ Address Resolution Protocol (reply)
 Hardware type: Ethernet (1)
 Protocol type: IPv4 (0x0800)
 Hardware size: 6
 Protocol size: 4
 Opcode: reply (2)
 Sender MAC address: Dell_c0:56:f0 (00:21:70:c0:56:f0)
 Sender IP address: 172.16.0.107
 Target MAC address: HewlettP_bf:91:ee (00:25:b3:bf:91:ee)
 Target IP address: 172.16.0.1

Addressing: routing to another LAN

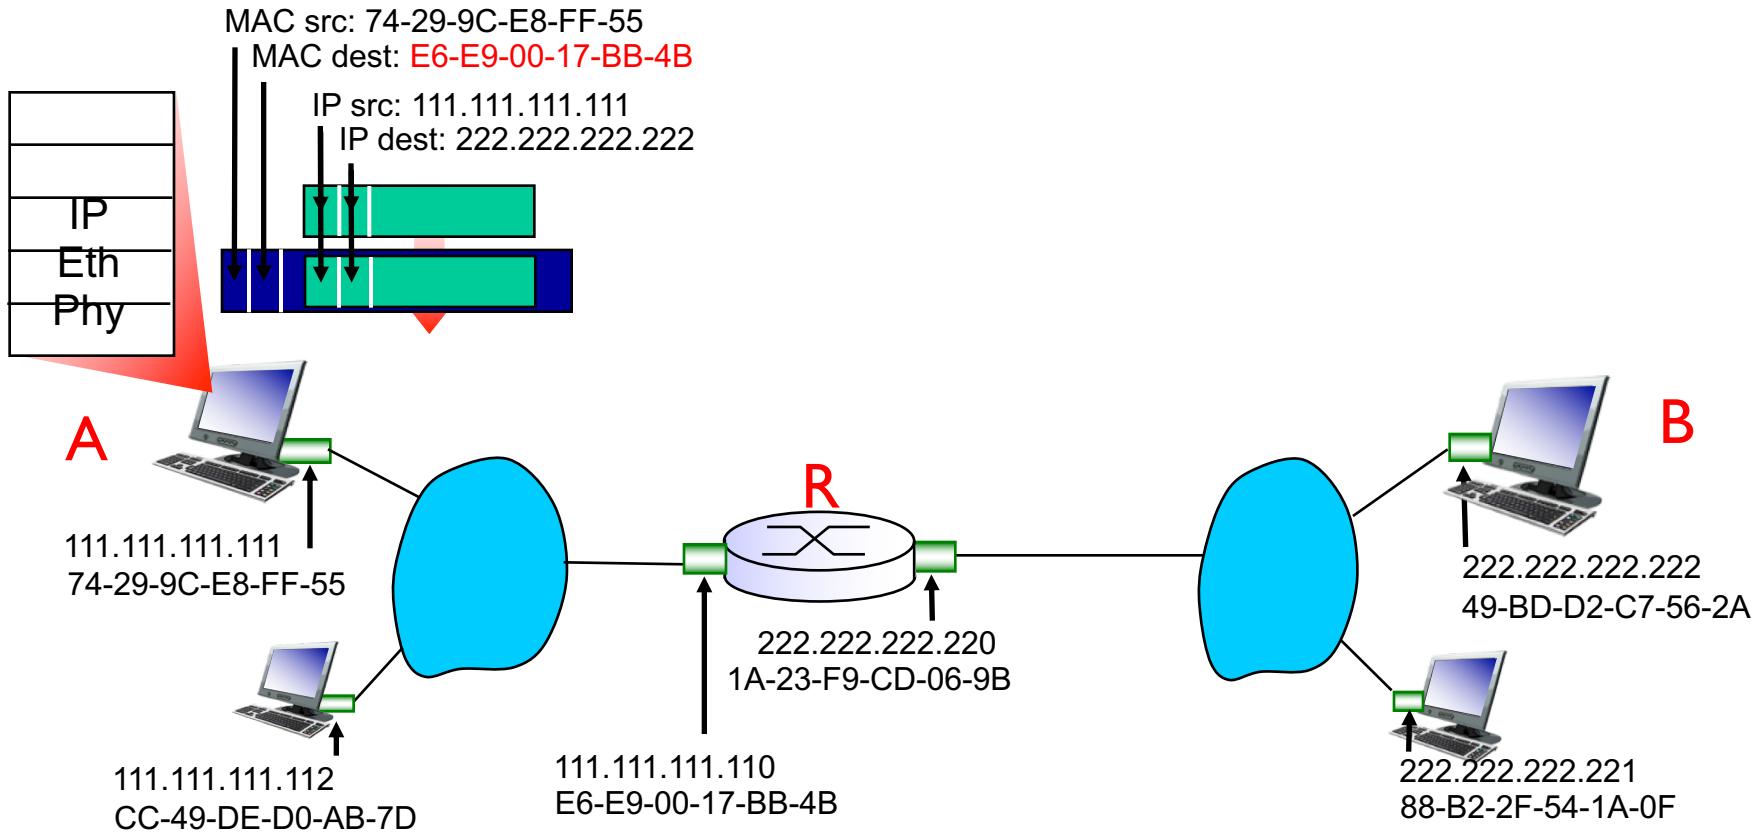
walkthrough: send datagram from A to B via R

- focus on addressing – at IP (datagram) and MAC layer (frame)
- assume A knows B's IP address
- assume A knows IP address of first hop router, R (how?)
- assume A knows R's MAC address (how?)



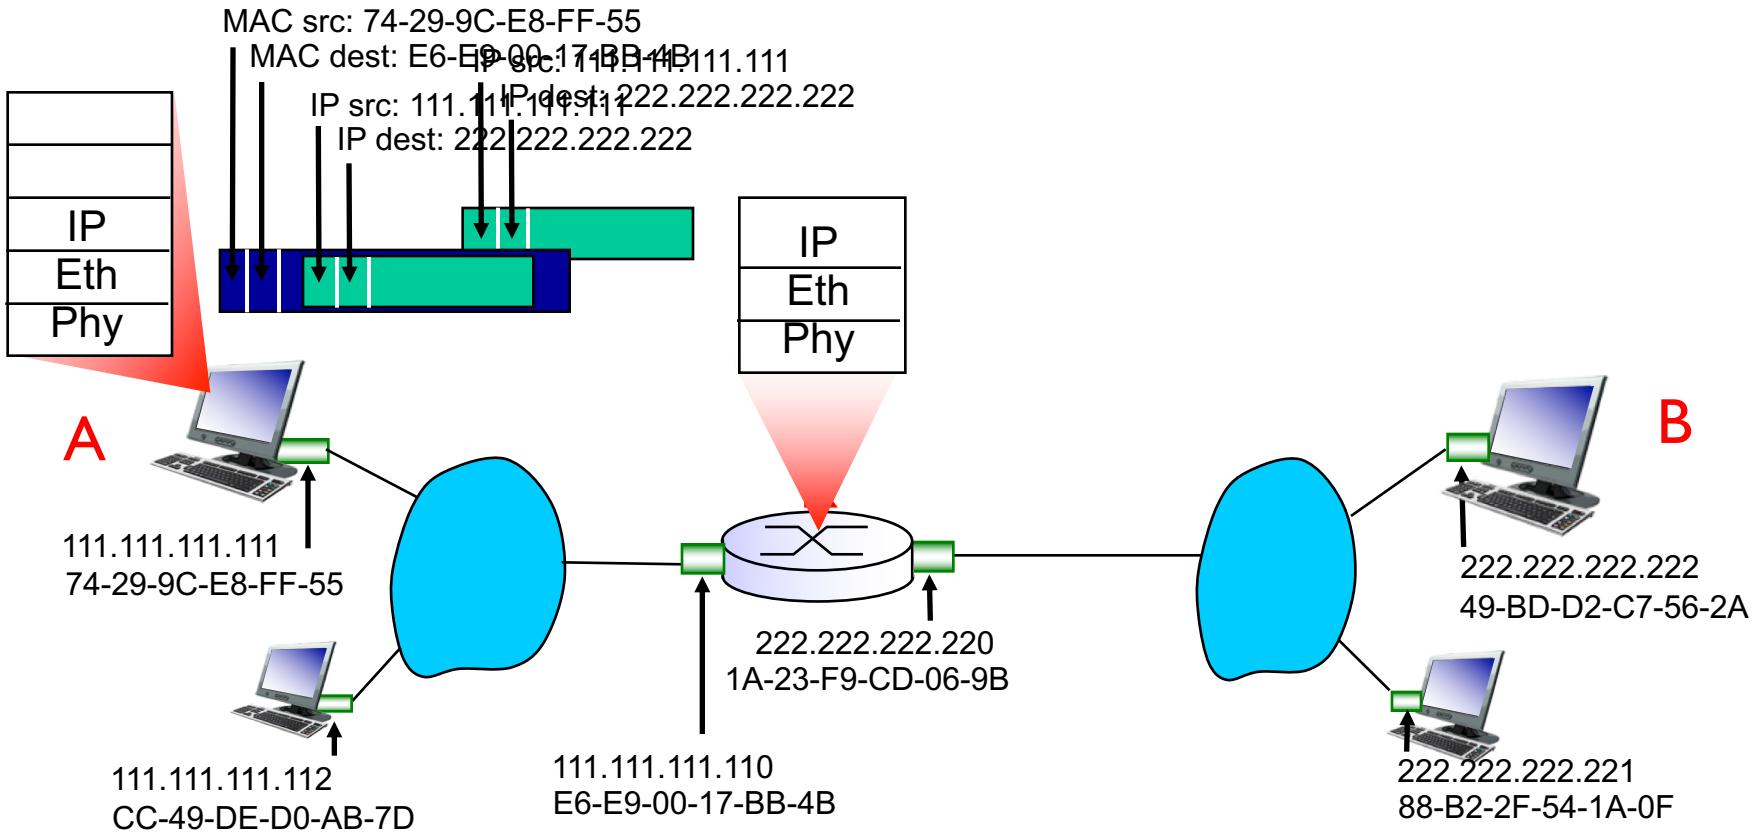
Addressing: routing to another LAN

- A creates IP datagram with IP source A, destination B
- A creates link-layer frame with R's MAC address as destination address, frame contains A-to-B IP datagram



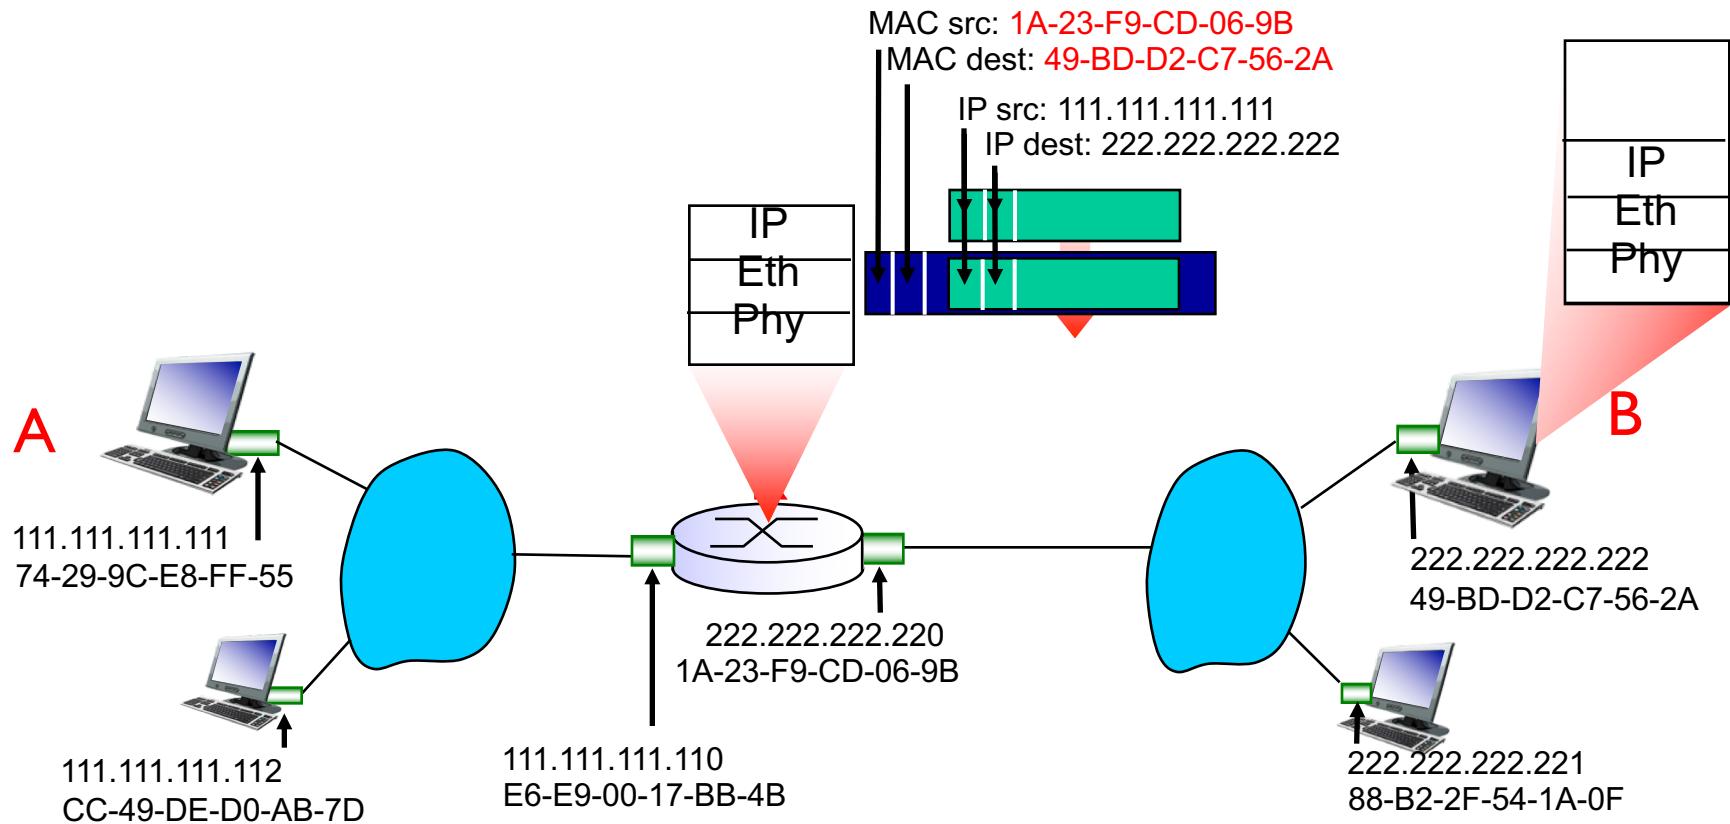
Addressing: routing to another LAN

- frame sent from A to R
- frame received at R, datagram removed, passed up to IP



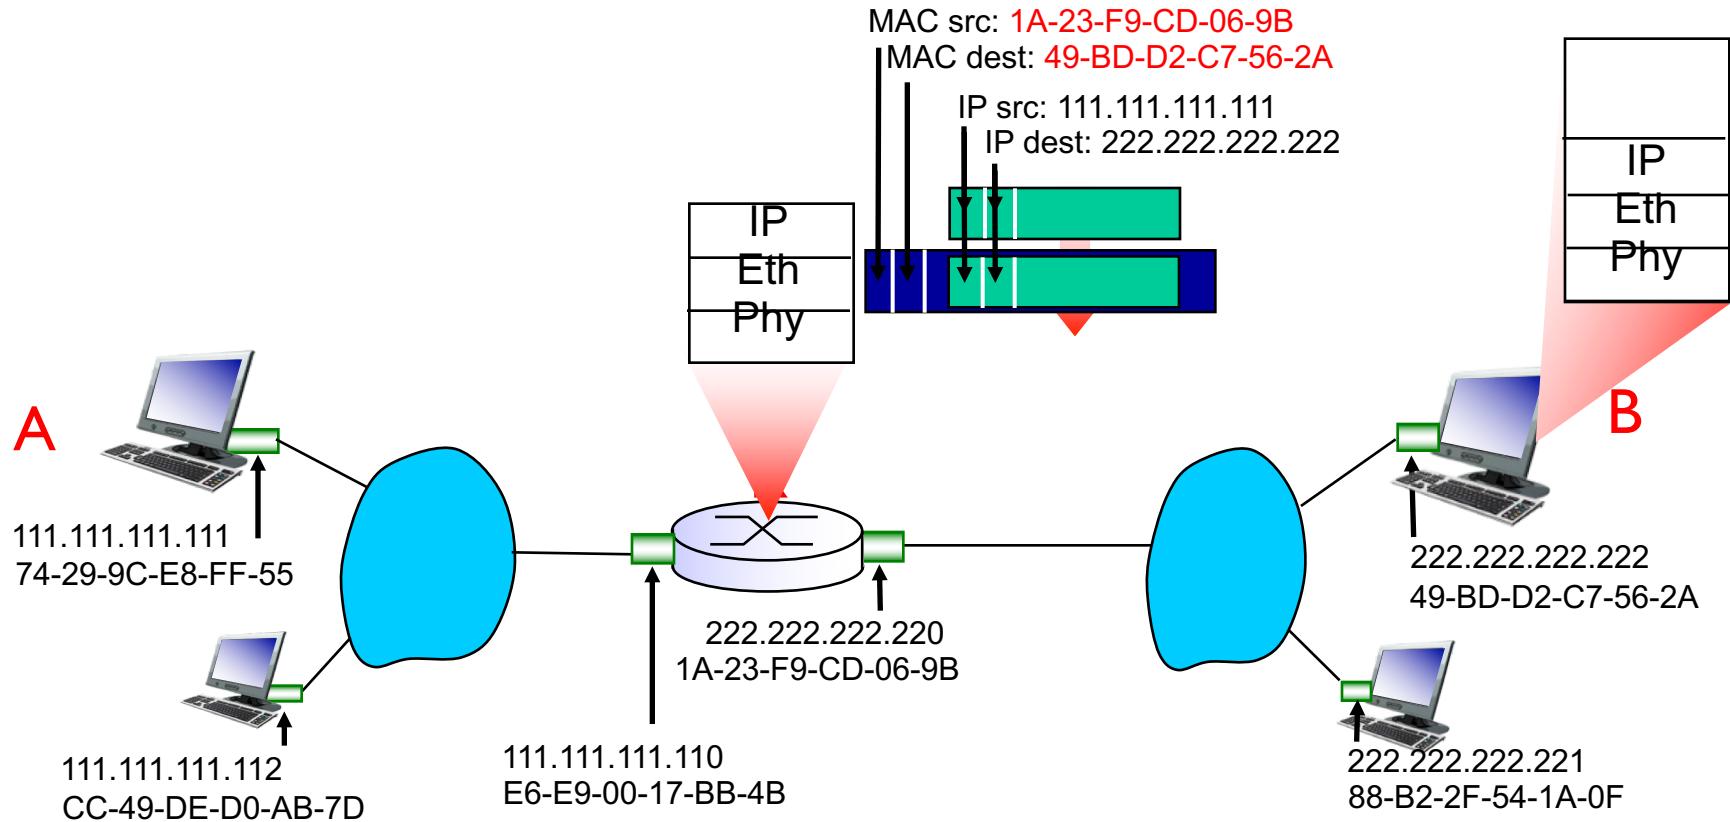
Addressing: routing to another LAN

- R forwards datagram with IP source A, destination B
- R creates link-layer frame with B's MAC address as destination address, frame contains A-to-B IP datagram



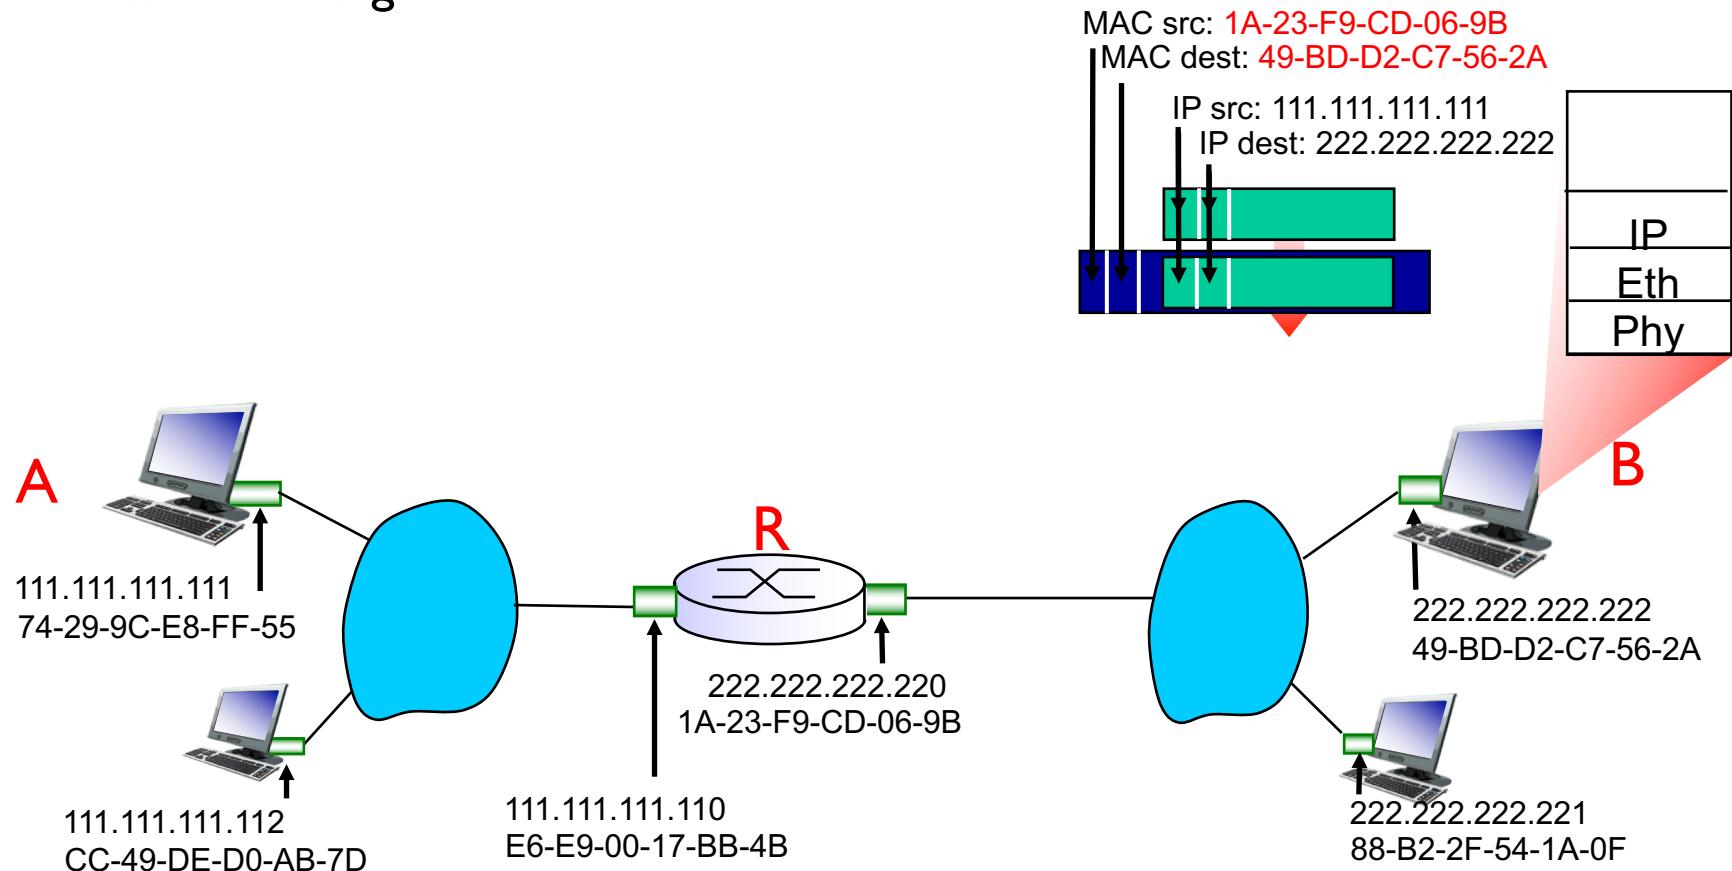
Addressing: routing to another LAN

- R forwards datagram with IP source A, destination B
- R creates link-layer frame with B's MAC address as destination address, frame contains A-to-B IP datagram



Addressing: routing to another LAN

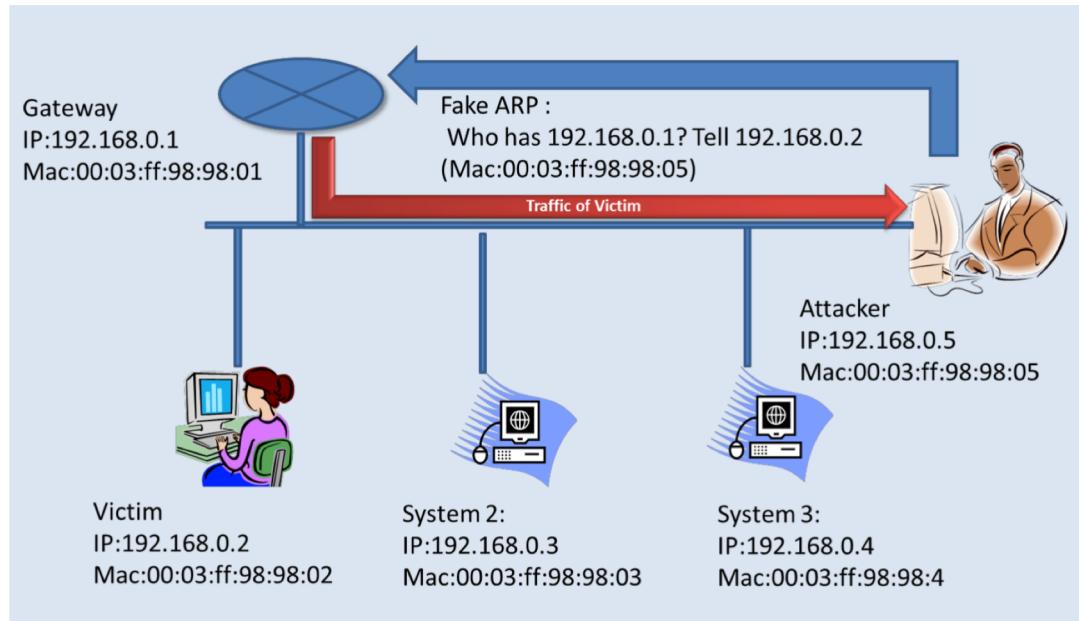
- R forwards datagram with IP source A, destination B
- R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

ARP Spoofing

- The ARP table is updated whenever an ARP response is received
- Requests are not tracked
- ARP announcements are not authenticated
- Machines trust each other
- A rogue machine can spoof other machines



ARP Poisoning (ARP Spoofing)

- According to the standard, almost all ARP implementations are stateless
- An arp cache updates every time that it receives an arp reply... even if it did not send any arp request!
- It is possible to “poison” an arp cache by sending **gratuitous arp replies**
- Using static entries solves the problem but it is almost impossible to manage!

Wireshark



- Wireshark is a packet sniffer and protocol analyzer
 - Captures and analyzes frames
 - Supports plugins
- Usually required to run with administrator privileges
- Setting the network interface in promiscuous mode captures traffic across the entire LAN segment and not just frames addressed to the machine
- Freely available on www.wireshark.org

(Untitled) - Wireshark

File Edit View Go Capture Analyze Statistics Help

← menu

main toolbar

Filter: Expression... Clear Apply

← filter toolbar

No. . Time Source Destination Protocol Info

1915	18.571194	212.97.59.91	128.148.36.11	UDP	Source port: 38662 Destination port: inovaport1
1916	18.587479	128.148.36.11	98.136.112.142	TCP	61219 > http [FIN, ACK] Seq=1 Ack=1 Win=16425 Len=0
1917	18.590200	128.148.36.11	212.97.59.91	UDP	Source port: inovaport1 Destination port: 38662
1918	18.591586	128.148.36.11	212.97.59.91	UDP	Source port: inovaport1 Destination port: 38662
1919	18.593191	212.97.59.91	128.148.36.11	UDP	Source port: 38662 Destination port: inovaport1
1920	18.602209	98.136.112.142	128.148.36.11	TCP	http > 61219 [ACK] Seq=1 Ack=1 Win=32850 Len=0
1921	18.604214	212.97.59.91	128.148.36.11	UDP	Source port: 38662 Destination port: inovaport1
1922	18.625996	128.148.36.11	212.97.59.91	UDP	Source port: inovaport1 Destination port: 38662
1923	18.626201	212.97.59.91	128.148.36.11	UDP	Source port: 38662 Destination port: inovaport1
1924	18.627287	128.148.36.11	212.97.59.91	UDP	Source port: inovaport1 Destination port: 38662
1925	18.648212	212.97.59.91	128.148.36.11	UDP	Source port: 38662 Destination port: inovaport1
1926	18.657224	128.148.36.11	212.97.59.91	UDP	Source port: inovaport1 Destination port: 38662
1927	18.670198	212.97.59.91	128.148.36.11	UDP	Source port: 38662 Destination port: inovaport1
1928	18.676199	98.136.112.142	128.148.36.11	TCP	http > 61219 [FIN, ACK] Seq=1 Ack=2 Win=32850 Len=0
1929	18.676289	128.148.36.11	98.136.112.142	TCP	61219 > http [ACK] Seq=2 Ack=2 Win=16425 Len=0
1930	18.686186	128.148.36.11	212.97.59.91	UDP	Source port: inovaport1 Destination port: 38662

Frame 1920 (60 bytes on wire, 60 bytes captured)

Ethernet II, Src: Micro-st_b2:d1:76 (00:0c:76:b2:d1:76), Dst: HewlettP_34:60:88 (00:22:64:34:60:88)

Destination: HewlettP_34:60:88 (00:22:64:34:60:88)

Source: Micro-st_b2:d1:76 (00:0c:76:b2:d1:76)

Type: IP (0x0800)

Trailer: 000000000000

Internet Protocol, Src: 98.136.112.142 (98.136.112.142), Dst: 128.148.36.11 (128.148.36.11)

Transmission Control Protocol, Src Port: http (80), Dst Port: 61219 (61219), Seq: 1, ACK: 2, Len: 0

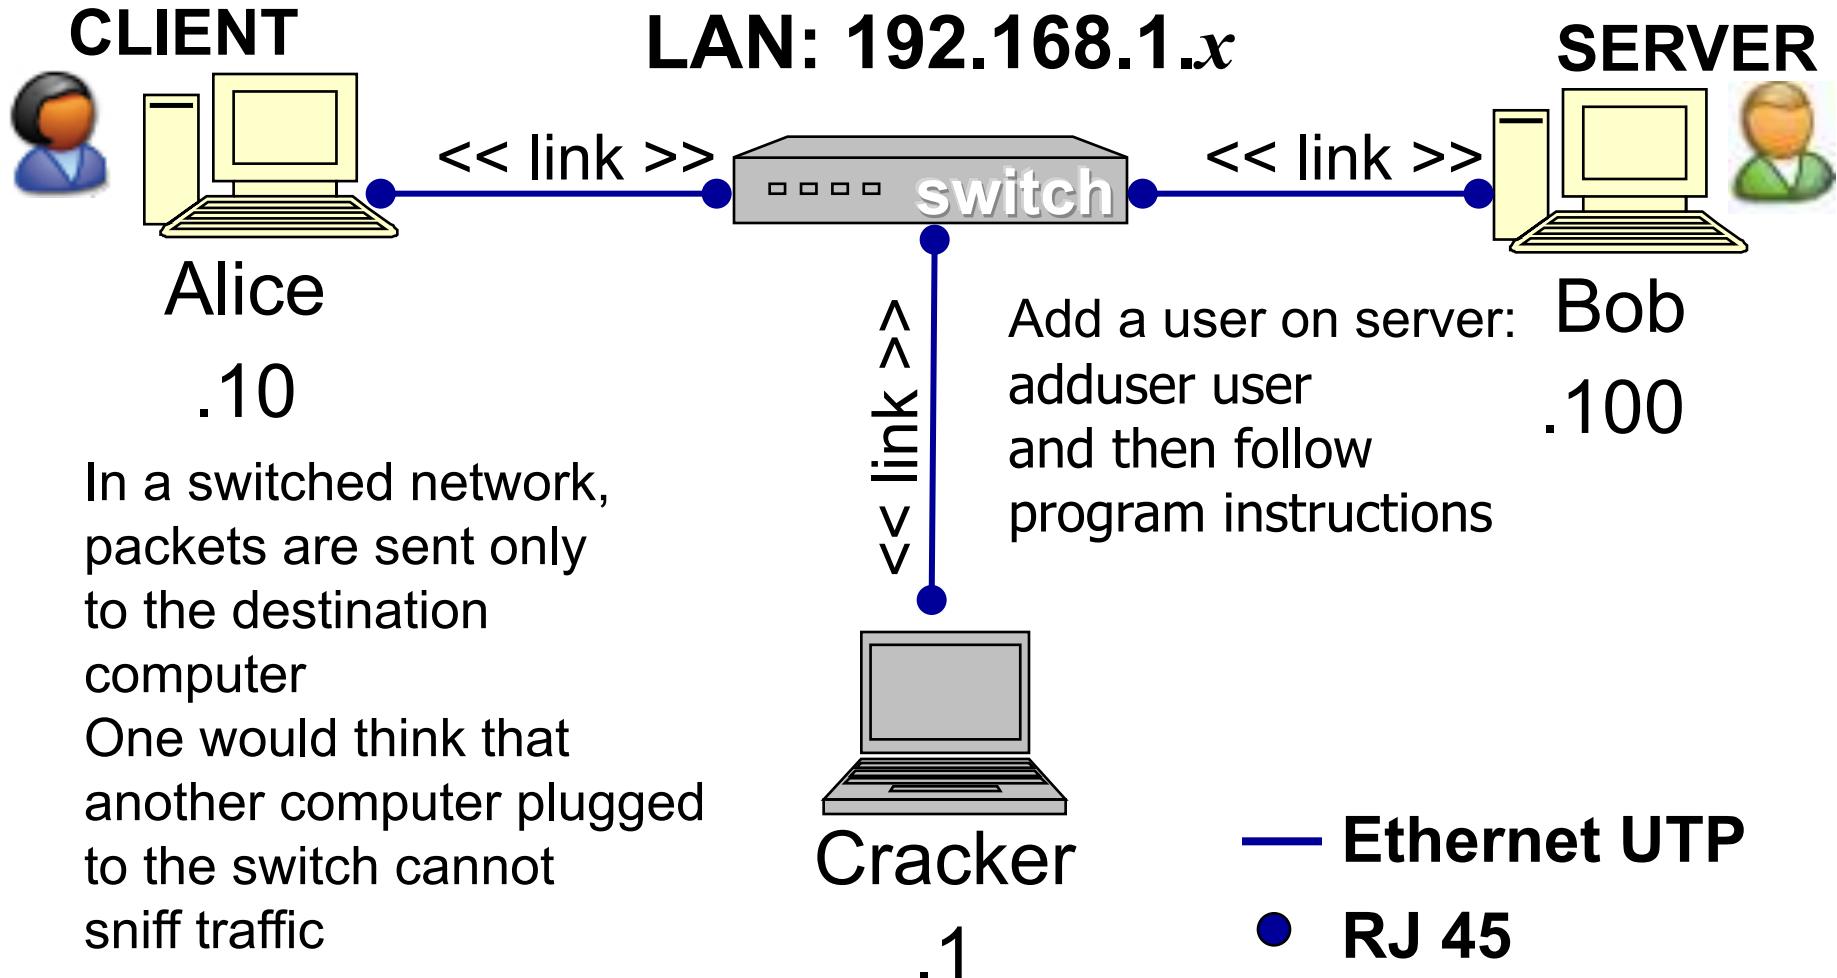
← packet details pane

← packet bytes pane

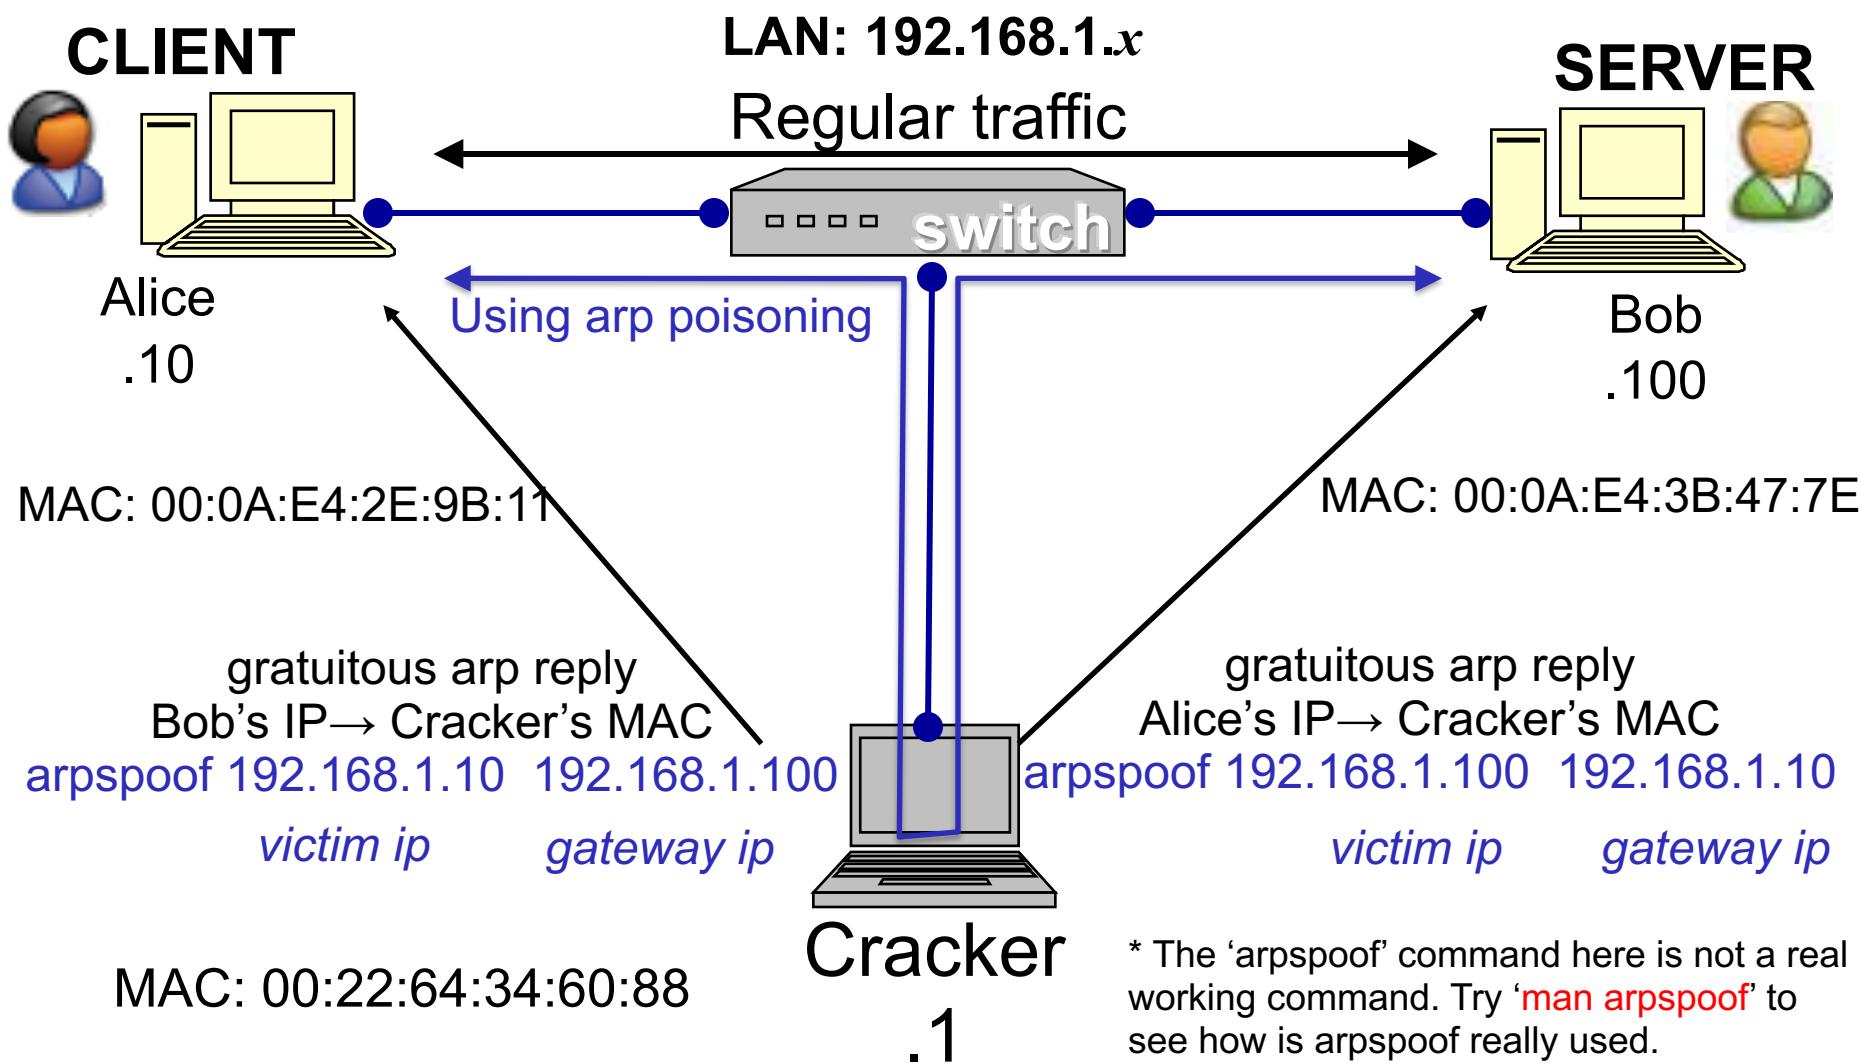
← status bar

27-Jan-21 Ethernet (eth), 20 bytes Packets: 2017 Displayed: 2017 Marked: 0 Dropped: 0 Networks Security 16

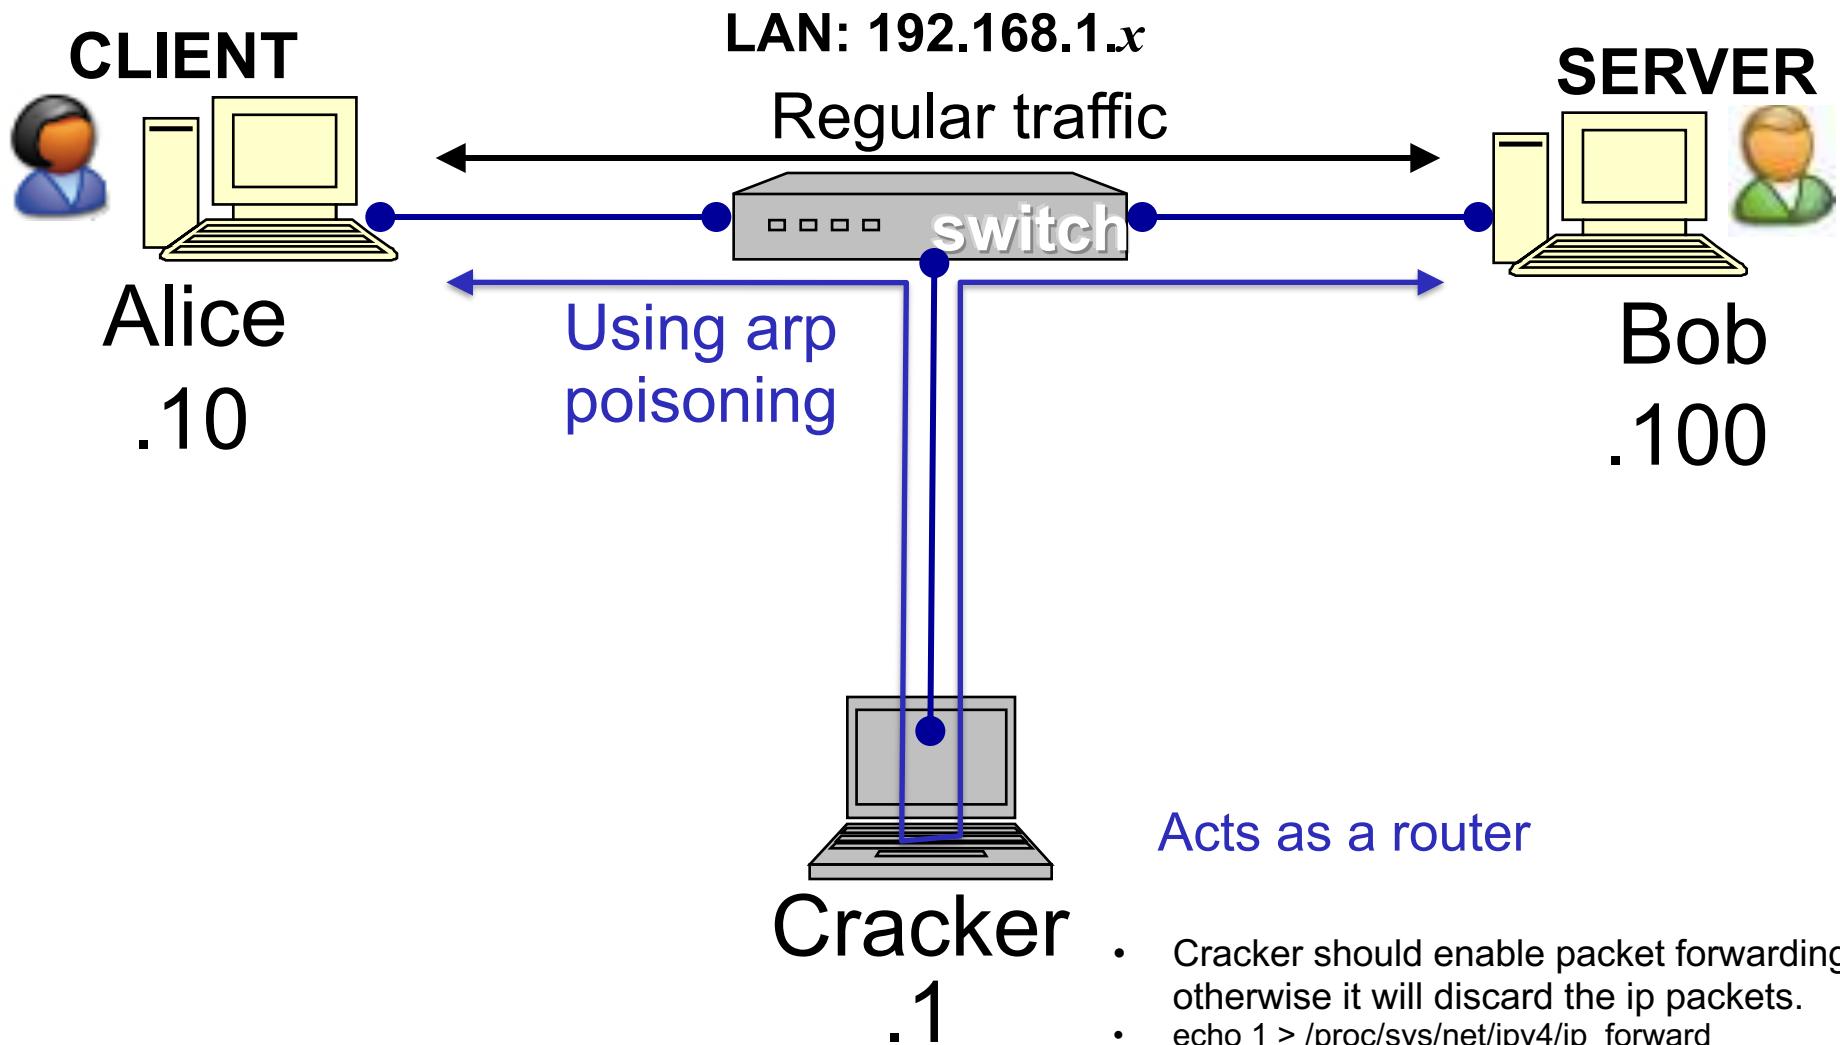
DEMO I: Configuration using Telnet



DEMO I: ARP Spoofing



DEMO I: catch telnet password



ARP Spoofing Detection

- Detected through duplicated MAC address

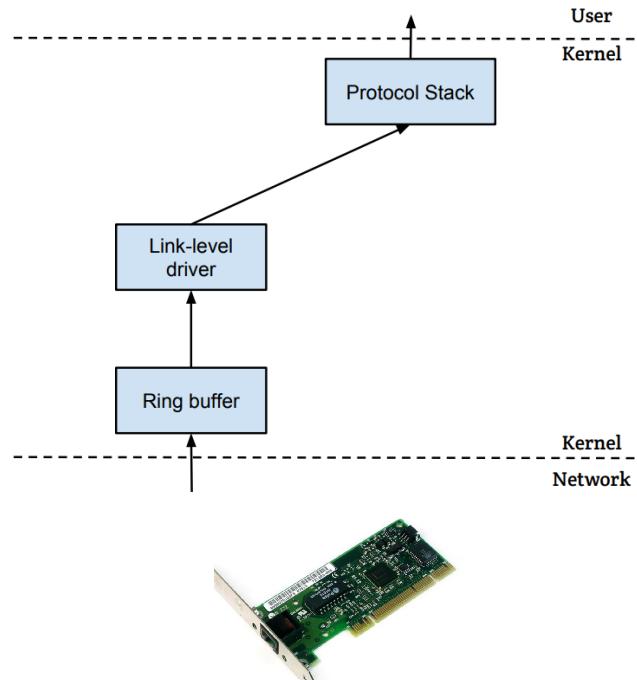
```
299 108.115463 Microsof_98:98:02  Microsof_98:98:01  ARP      60 who has 192.168.0.1? Tell 192.168.0.3 (duplicate use of 192.168.0.3 detected!)
300 108.115463 Microsof_98:98:01  Microsof_98:98:02  ARP      42 192.168.0.1 is at 00:03:ff:98:98:01 (duplicate use of 192.168.0.3 detected!)
301 108.115463 Microsof_98:98:02  Microsof_98:98:01  ARP      60 192.168.0.3 is at 00:03:ff:98:98:02
302 138.198720 Microsof_98:98:02  Microsof_98:98:01  ARP      60 192.168.0.3 is at 00:03:ff:98:98:02
319 168.231906 Microsof_98:98:02  Microsof_98:98:01  ARP      60 192.168.0.3 is at 00:03:ff:98:98:02

Frame 299: 60 bytes on wire (480 bits), 60 bytes captured (480 bits)
Ethernet II, Src: Microsof_98:98:02 (00:03:ff:98:98:02), Dst: Microsof_98:98:01 (00:03:ff:98:98:01)
[Duplicate IP address detected for 192.168.0.3 (00:03:ff:98:98:02) - also in use by 00:03:ff:98:98:03 (frame 23)]
Address Resolution Protocol (request)
Hardware type: Ethernet (1)
Protocol type: IP (0x0800)
Hardware size: 6
Protocol size: 4
Opcode: request (1)
Sender MAC address: Microsof_98:98:02 (00:03:ff:98:98:02)
Sender IP address: 192.168.0.3 (192.168.0.3)
Target MAC address: 00:00:00_00:00:00 (00:00:00:00:00:00)
Target IP address: 192.168.0.1 (192.168.0.1)
```

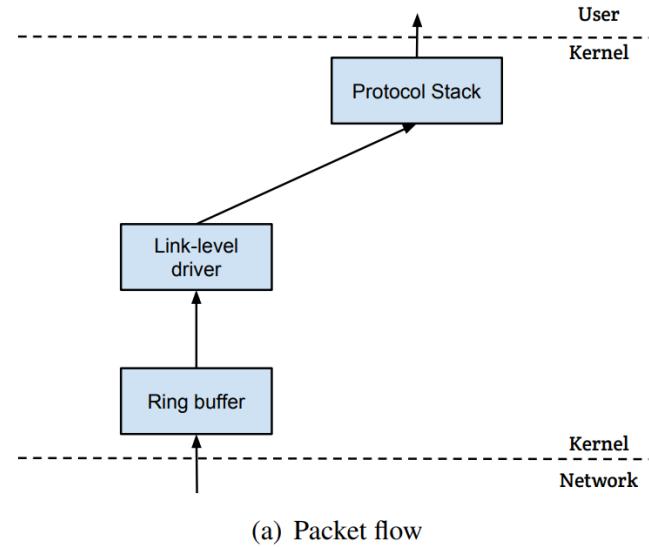
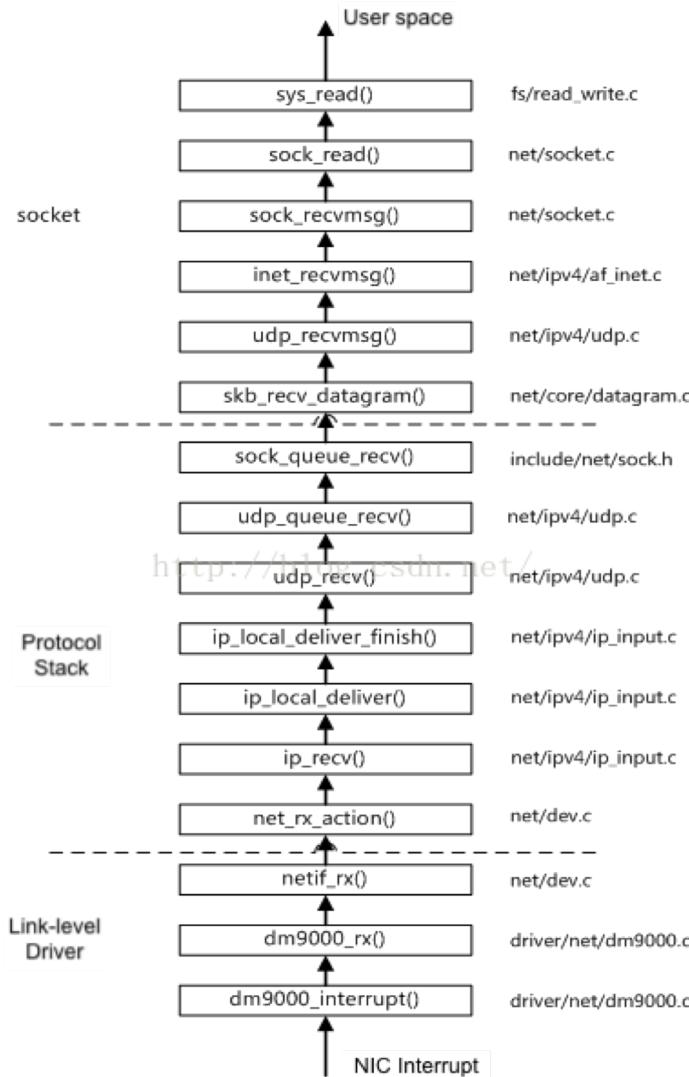
Packet Sniffing and Spoofing

How Packets Are Received

- NIC (Network Interface Card) is a physical or logical link between a machine and a network
- Each NIC has a MAC address
- Every NIC on the network will hear all the frames on the wire
- NIC checks the destination address for every packet, if the address matches the cards MAC address, it is further copied into a buffer in the kernel
 - Copy through direct memory access (DMA)
 - Then NIC triggers an interrupt



What indeed happens inside a Linux kernel



(a) Packet flow

Q: where are ARP and ICMP processed?

Promiscuous Mode

- The frames that are not destined to a given NIC are discarded
 - i.e., if they do not match the destination **MAC address**
- When operating in promiscuous mode, NIC passes every frame received from the **network** to the **kernel**
 - promiscuous mode is a **NIC** mode!
- If a sniffer program is registered with the kernel, it will be able to see all the packets
 - **Q:** where are the packets now?
- In Wi-Fi, it is called Monitor Mode

BSD Packet Filter (BPF)

```
struct sock_filter code[] = {
{ 0x28, 0, 0, 0x0000000c }, { 0x15, 0, 8, 0x000086dd },
{ 0x30, 0, 0, 0x00000014 }, { 0x15, 2, 0, 0x00000084 },
{ 0x15, 1, 0, 0x00000006 }, { 0x15, 0, 17, 0x00000011 },
{ 0x28, 0, 0, 0x00000036 }, { 0x15, 14, 0, 0x00000016 },
{ 0x28, 0, 0, 0x00000038 }, { 0x15, 12, 13, 0x00000016 },
{ 0x15, 0, 12, 0x000000800 }, { 0x30, 0, 0, 0x00000017 },
{ 0x15, 2, 0, 0x00000084 }, { 0x15, 1, 0, 0x00000006 },
{ 0x15, 0, 8, 0x00000011 }, { 0x28, 0, 0, 0x00000014 },
{ 0x45, 6, 0, 0x00001fff }, { 0xb1, 0, 0, 0x0000000e },
{ 0x48, 0, 0, 0x0000000e }, { 0x15, 2, 0, 0x00000016 },
{ 0x48, 0, 0, 0x00000010 }, { 0x15, 0, 1, 0x00000016 },
{ 0x06, 0, 0, 0x0000ffff }, { 0x06, 0, 0, 0x00000000 },
};

struct sock_fprog bpf = {
    .len = ARRAY_SIZE(code),
    .filter = code,
};
```

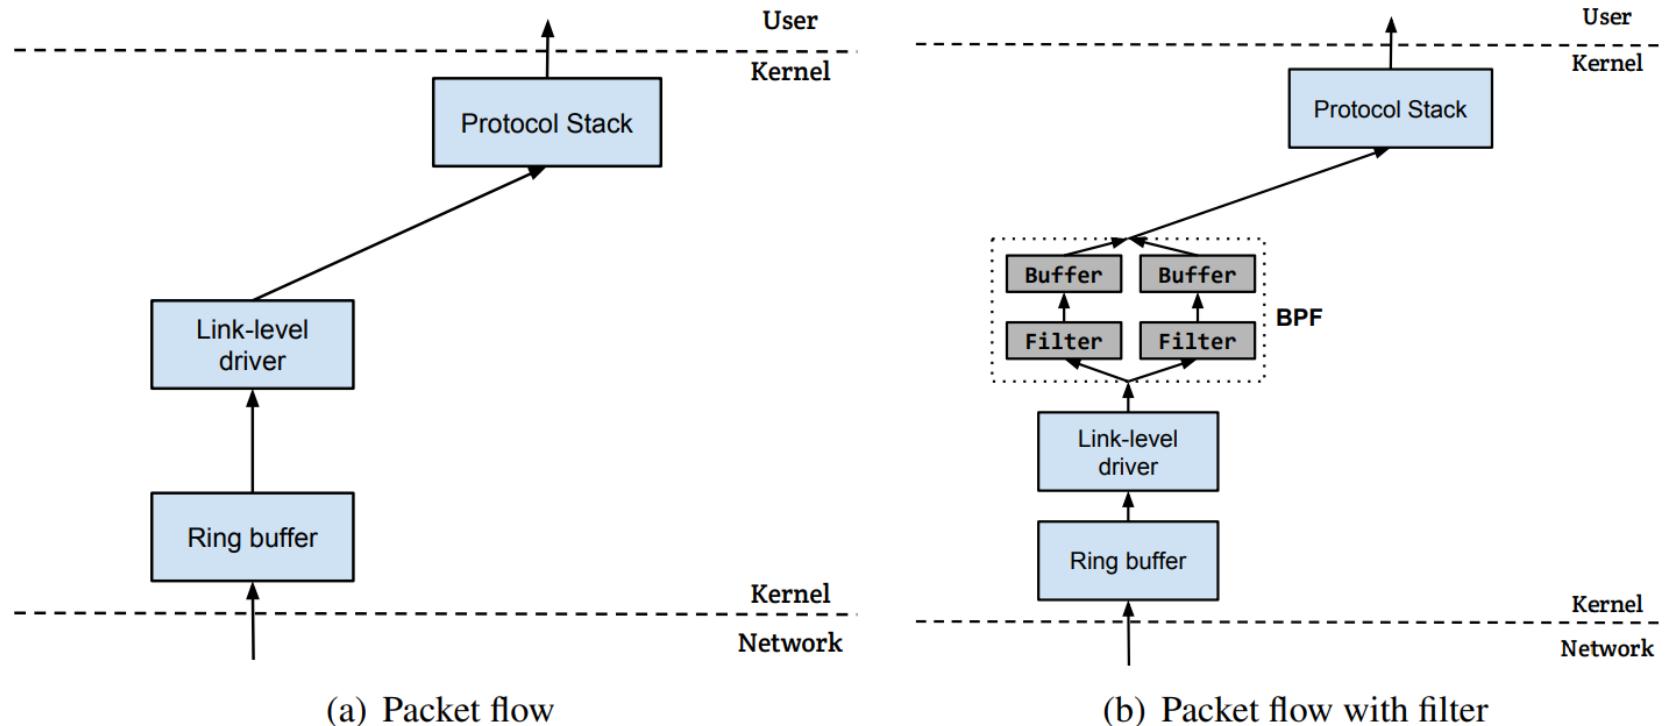
- **BPF allows a user-program to attach a filter to the socket, which tells the kernel to discard unwanted packets.**
- **An example of the compiled BPF code is shown here.**

BSD Packet Filter (BPF)

```
setsockopt(sock, SOL_SOCKET, SO_ATTACH_FILTER, &bpf, sizeof(bpf))
```

- A compiled BPF pseudo-code can be attached to a socket through `setsockopt()`
- When a packet is received by kernel, BPF will be invoked
- An accepted packet is pushed up the protocol stack. See the diagram on the following slide.

Packet Flow With/Without Filters



Packet Sniffing

Packet sniffing describes the process of capturing live data as they flow across a network

Let's first see how computers receive packets.

Receiving Packets Using Socket

Create the socket

```
// Step ①
int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);

// Step ②
memset((char *) &server, 0, sizeof(server));
server.sin_family = AF_INET;
server.sin_addr.s_addr = htonl(INADDR_ANY);
server.sin_port = htons(9090);

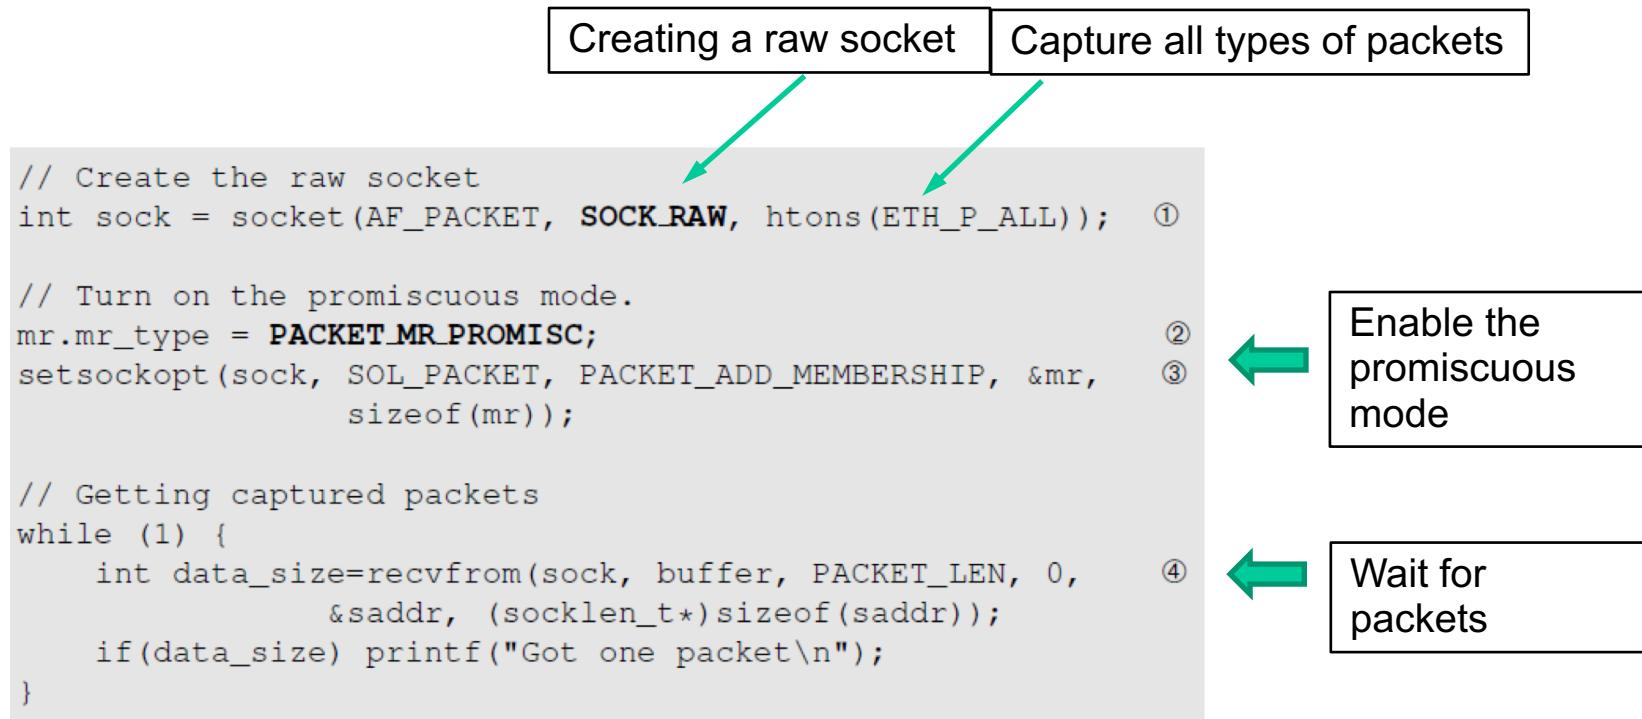
if (bind(sock, (struct sockaddr *) &server, sizeof(server)) < 0)
    error("ERROR on binding");
```

Provide information about server

```
// Step ③
while (1) {
    bzero(buf, 1500);
    recvfrom(sock, buf, 1500-1, 0,
             (struct sockaddr *) &client, &clientlen);
    printf("%s\n", buf);
}
```

Receive packets

Receiving Packets Using Raw Socket



Limitation of the Approach

- This program is not portable across different operating systems.
- Setting filters is not easy.
- The program does not explore any optimization to improve performance.
- The PCAP library was thus created.
 - It still uses raw sockets internally, but its API is standard across all platforms. OS specifics are hidden by PCAP's implementation.
 - Allows programmers to specify filtering rules using human readable Boolean expressions.

Packet Sniffing Using the pcap API

```
char filter_exp[] = "ip proto icmp";
```

```
// Step 1: Open live pcap session on NIC with name eth3  
handle = pcap_open_live("eth3", BUFSIZ, 1, 1000, errbuf); ①
```

Filter

```
// Step 2: Compile filter_exp into BPF psuedo-code  
pcap_compile(handle, &fp, filter_exp, 0, net); ②  
pcap_setfilter(handle, &fp); ③
```

```
// Step 3: Capture packets  
pcap_loop(handle, -1, got_packet, NULL); ④
```

Initialize a raw socket, set the network device into promiscuous mode.

Invoke this function for every captured packet

```
void got_packet(u_char *args, const struct pcap_pkthdr *header,  
                const u_char *packet)  
{  
    printf("Got a packet\n");  
}
```

Processing Captured Packet: Ethernet Header

```
/* Ethernet header */
struct ethheader {
    u_char ether_dhost[ETHER_ADDR_LEN]; /* destination host address */
    u_char ether_shost[ETHER_ADDR_LEN]; /* source host address */
    u_short ether_type;                 /* IP? ARP? RARP? etc */
};

void got_packet(u_char *args, const struct pcap_pkthdr *header,
                const u_char *packet)
{
    struct ethheader *eth = (struct ethheader *)packet;
    if ( ntohs(eth->ether_type) == 0x0800) { ... } // IP packet
    ...
}
```

The **packet** argument contains a copy of the packet, including the Ethernet header. We typecast it to the Ethernet header structure.

Now we can access the field of the structure

Processing Captured Packet: IP Header

```
void got_packet(u_char *args, const struct pcap_pkthdr *header,
                 const u_char *packet)
{
    struct ethheader *eth = (struct ethheader *)packet;

    if ( ntohs(eth->ether_type) == 0x0800) { // 0x0800 is IP type
        struct ipheader * ip = (struct ipheader *)
            (packet + sizeof(struct ethheader)); ①

        printf("      From: %s\n", inet_ntoa(ip->iph_sourceip)); ②
        printf("      To: %s\n", inet_ntoa(ip->iph_destip)); ③

        /* determine protocol */
        switch(ip->iph_protocol) {
            case IPPROTO_TCP:
                printf("      Protocol: TCP\n");
                return;
            case IPPROTO_UDP:
                printf("      Protocol: UDP\n");
                return;
        }
    }
}
```

Find where the IP header starts, and typecast it to the IP Header structure.

Now we can easily access the fields in the IP header.

Further Processing Captured Packet

- If we want to further process the packet, such as printing out the header of the TCP, UDP and ICMP, we can use the similar technique.
 - We move the pointer to the beginning of the next header and type-cast
 - We need to use the header length field in the IP header to calculate the actual size of the IP header
- In the following example, if we know the next header is ICMP, we can get a pointer to the ICMP part by doing the following:

```
int ip_header_len = ip->iph_ihl * 4;  
u_char *icmp = (struct icmpheader *)  
    (packet + sizeof(struct ethheader) + ip_header_len);
```

Packet Spoofing

- When some critical information in the packet is forged, we refer to it as packet spoofing.
- Many network attacks rely on packet spoofing.
- Let's see how to send packets without spoofing.
- Code available on eDimension or <http://github.com/kevin-w-du/BookCode>

Sending Packets Without Spoofing

```
void main()
{
    struct sockaddr_in dest_info;
    char *data = "UDP message\n";

    // Step 1: Create a network socket
    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);

    // Step 2: Provide information about destination.
    memset((char *) &dest_info, 0, sizeof(dest_info));
    dest_info.sin_family = AF_INET;
    dest_info.sin_addr.s_addr = inet_addr("10.0.2.5");
    dest_info.sin_port = htons(9090);

    // Step 3: Send out the packet.
    sendto(sock, data, strlen(data), 0,
           (struct sockaddr *)&dest_info, sizeof(dest_info));
    close(sock);
}
```

Testing: Use the netcat (nc) command to run a UDP server on 10.0.2.5. We then run the program on the left from another machine. We can see that the message has been delivered to the server machine:



```
seed@Server(10.0.2.5):$ nc -lув 9090
Connection from 10.0.2.6 port 9090 [udp/*] accepted
UDP message
```

Spoofing Packets Using Raw Sockets

There are two major steps in packet spoofing:

- Constructing the packet
- Sending the packet out

Spoofing Packets Using Raw Sockets

```
/*
 * Given an IP packet, send it out using a raw socket.
 */
void send_raw_ip_packet(struct ipheader* ip)
{
    struct sockaddr_in dest_info;
    int enable = 1;

    // Step 1: Create a raw network socket.
    int sock = socket(AF_INET, SOCK_RAW, IPPROTO_RAW);

    // Step 2: Set socket option.
    setsockopt(sock, IPPROTO_IP, IP_HDRINCL,
               &enable, sizeof(enable));

    // Step 3: Provide needed information about destination.
    dest_info.sin_family = AF_INET;
    dest_info.sin_addr = ip->iph_destip;

    // Step 4: Send the packet out.
    sendto(sock, ip, ntohs(ip->iph_len), 0,
           (struct sockaddr *)&dest_info, sizeof(dest_info));
    close(sock);
}
```

We use `setsockopt()` to enable `IP_HDRINCL` on the socket.

For raw socket programming, since the destination information is already included in the provided IP header, we do not need to fill all the fields

Since the socket type is raw socket, the system will send out the IP packet as is.

Spoofing Packets: Constructing the Packet

Fill in the ICMP Header

```
char buffer[1500];  
  
memset(buffer, 0, 1500);  
  
/*********************************************  
 Step 1: Fill in the ICMP header.  
 ****/  
struct icmpheader *icmp = (struct icmpheader *)  
    (buffer + sizeof(struct ipheader));  
icmp->icmp_type = 8; //ICMP Type: 8 is request, 0 is reply.  
  
// Calculate the checksum for integrity  
icmp->icmp_chks = 0;  
icmp->icmp_chks = in_cksum((unsigned short *)icmp,  
                           sizeof(struct icmpheader));
```

Find the starting point of the ICMP header, and typecast it to the ICMP structure

Fill in the ICMP header fields

Spoofing Packets: Constructing the Packet

Fill in the IP Header

```
/*********************  
 Step 2: Fill in the IP header.  
*****  
struct ipheader *ip = (struct ipheader *) buffer;  
ip->iph_ver = 4;  
ip->iph_ihl = 5;  
ip->iph_ttl = 20;  
ip->iph_sourceip.s_addr = inet_addr("1.2.3.4");  
ip->iph_destip.s_addr = inet_addr("10.0.2.5");  
ip->iph_protocol = IPPROTO_ICMP;  
ip->iph_len = htons(sizeof(struct ipheader) +  
                     sizeof(struct icmpheader));
```

Typecast the buffer
to the IP structure

Fill in the IP header
fields

Finally, send out the packet

```
send_raw_ip_packet (ip);
```

Spoofing UDP Packets

```
memset(buffer, 0, 1500);
struct ipheader *ip = (struct ipheader *) buffer;
struct udpheader *udp = (struct udpheader *) (buffer +
                                              sizeof(struct ipheader));

/*****************
 Step 1: Fill in the UDP data field.
 *****/
char *data = buffer + sizeof(struct ipheader) +
             sizeof(struct udpheader);
const char *msg = "Hello Server!\n";
int data_len = strlen(msg);
strncpy (data, msg, data_len);

/*****************
 Step 2: Fill in the UDP header.
 *****/
udp->udp_sport = htons(12345);
udp->udp_dport = htons(9090);
udp->udp_ulen = htons(sizeof(struct udpheader) + data_len);
udp->udp_sum = 0; /* Many OSes ignore this field, so we do not
                     calculate it. */
```

↳ Constructing UDP packets is similar, except that we need to include the payload data now.

Spoofing UDP Packets (continued)

```
*****  
Step 3: Fill in the IP header.  
*****  
..... /* Code omitted here; same as that in Listing 12.6 */  
ip->iph_protocol = IPPROTO_UDP; // The value is 17.  
ip->iph_len = htons(sizeof(struct ipheader) +  
                     sizeof(struct udphdr) + data_len);
```

Testing: Use the nc command to run a UDP server on 10.0.2.5. We then spoof a UDP packet from another machine. We can see that the spoofed UDP packet was received by the server machine.

```
seed@Server(10.0.2.5):$ nc -lув 9090  
Connection from 1.2.3.4 port 9090 [udp/*] accepted  
Hello Server!
```

Sniffing and Then Spoofing

- In many situations, we need to capture packets first, and then spoof a response based on the captured packets.
- Procedure (using UDP as example)
 - Use PCAP API to capture the packets of interests
 - Make a copy from the captured packet
 - Replace the UDP data field with a new message and swap the source and destination fields
 - Send out the spoofed reply

UDP Packet

```
void spoof_reply(struct ipheader* ip)
{
    const char buffer[1500];
    int ip_header_len = ip->iph_ihl * 4;
    struct udpheader* udp = (struct udpheader *) ((u_char *)ip +
                                                ip_header_len);

    if ( ntohs(udp->udp_dport) != 9999) {
        // Only spoof UDP packet with destination port 9999
        return;
    }

    // Step 1: Make a copy from the original packet
    memset((char*)buffer, 0, 1500);
    memcpy((char*)buffer, ip, ntohs(ip->iph_len));
    struct ipheader * newip = (struct ipheader *) buffer;
    struct udpheader * newudp = (struct udpheader *) (buffer +
                                                       ip_header_len);
    char *data = (char *)newudp + sizeof(struct udpheader);

    // Step 2: Construct the UDP payload, keep track of payload size
    const char *msg = "This is a spoofed reply!\n";
    int data_len = strlen(msg);
    strncpy (data, msg, data_len);
```

UDP Packet (Continued)

```
// Step 3: Construct the UDP Header
newudp->udp_sport = udp->udp_dport;
newudp->udp_dport = udp->udp_sport;
newudp->udp_ulen = htons(sizeof(struct udpheader) + data_len);
newudp->udp_sum = 0;

// Step 4: Construct the IP header (no change for other fields)
newip->iph_sourceip = ip->iph_destip;
newip->iph_destip = ip->iph_sourceip;
newip->iph_ttl = 50; // Rest the TTL field
newip->iph_len = htons(sizeof(struct ipheader) +
                      sizeof(struct udpheader) + data_len);

// Step 5: Send out the spoofed IP packet
send_raw_ip_packet(newip);
}
```

Packing Sniffing Using Scapy

```
#!/usr/bin/python3
from scapy.all import *

print("SNIFFING PACKETS.....")

def print_pkt(pkt) : ①
    print("Source IP:", pkt[IP].src)
    print("Destination IP:", pkt[IP].dst)
    print("Protocol:", pkt[IP].proto)
    print("\n")

pkt = sniff(filter='icmp', prn=print_pkt) ②
```

Spoofing ICMP & UDP Using Scapy

```
#!/usr/bin/python3
from scapy.all import *

print("SENDING SPOOFED ICMP PACKET.....")
ip = IP(src="1.2.3.4", dst="93.184.216.34")    ①
icmp = ICMP()                                     ②
pkt = ip/icmp                                      ③
pkt.show()
send(pkt,verbose=0)                                ④
```

```
#!/usr/bin/python3
from scapy.all import *

print("SENDING SPOOFED UDP PACKET.....")
ip = IP(src="1.2.3.4", dst="10.0.2.69") # IP Layer
udp = UDP(sport=8888, dport=9090)        # UDP Layer
data = "Hello UDP!\n"                      # Payload
pkt = ip/udp/data                         # Construct the complete packet
pkt.show()
send(pkt,verbose=0)
```

Sniffing and Then Spoofing Using Scapy

```
#!/usr/bin/python3
from scapy.all import *

def spoof_pkt(pkt):
    if ICMP in pkt and pkt[ICMP].type == 8:
        print("Original Packet.....")
        print("Source IP : ", pkt[IP].src)
        print("Destination IP :", pkt[IP].dst)

        ip = IP(src=pkt[IP].dst, dst=pkt[IP].src, ihl=pkt[IP].ihl)
        icmp = ICMP(type=0, id=pkt[ICMP].id, seq=pkt[ICMP].seq)
        data = pkt[Raw].load
        newpkt = ip/icmp/data

        print("Spoofed Packet.....")
        print("Source IP : ", newpkt[IP].src)
        print("Destination IP :", newpkt[IP].dst)
        send(newpkt,verbose=0)

pkt = sniff(filter='icmp and src host 10.0.2.69',prn=spoof_pkt)
```

Packet Spoofing: Scapy v.s C

- Python + Scapy
 - Pros: constructing packets is very simple
 - Cons: much slower than C code
- C Program (using raw socket)
 - Pros: much faster
 - Cons: constructing packets is complicated
- Hybrid Approach
 - Using Scapy to construct packets
 - Using C to slightly modify packets and then send packets

Endianness

- Endianness: a term that refers to the order in which a given multi-byte data item is stored in memory.
 - **LittleEndian**: store the most significant byte of data at the highest address
 - **BigEndian**: store the most significant byte of data at the lowest address



Endianness In Network Communication

- Computers with different byte orders will “misunderstand” each other.
 - Solution: agree upon a common order for communication
 - This is called “network order”, which is the same as big endian order
- All computers need to convert data between “host order” and “network order” .

Macro	Description
htons ()	Convert unsigned short integer from host order to network order.
htonl ()	Convert unsigned integer from host order to network order.
ntohs ()	Convert unsigned short integer from network order to host order.
ntohl ()	Convert unsigned integer from network order to host order.

Summary

- Packet sniffing
 - Using raw socket
 - Using PCAP APIs
- Packet spoofing using raw socket
- Sniffing and the spoofing
- Endianness